

Optimal Policies



Frederik Mallmann-Trenn
6CCS3AIN

How utilities are calculated

- So far we have assumed that utilities are summed along a run.
 - Not the only way.
- In general we need to compute $U_r([s_0, s_1, \dots, s_n])$ for general $U_r(\cdot)$. That is, the utility of a run.
- Before $U_r(\cdot)$ was just the sum of rewards in every state.
- Can consider **finite** and **infinite** horizons.
 - Is it “game over” at some point?
- Turns out that infinite horizons are mostly easier to deal with.
 - That is what we will use.

How utilities are calculated

- Also have to consider whether **utilities** are **stationary** or **non-stationary**.
 - Think of: does the same state always have the same value?
 - E.g., in Pacman when you pick up a fruit, there is a large reward for that tile. That changes after you picked up the fruit.
- Example:
 - Normally we prefer one state to another.
 - Passing the AI module to failing it
 - In this case when the exam is, today or next week, is irrelevant.
- We assume utilities are **stationary**.

But are they?

- Not clear that utilities are always stationary.



- In truth, I don't always most want to eat cherry pie.
- Despite this, we will assume that utilities are stationary.

How utilities are calculated

- With stationary utilities, there are two ways to establish $U_r([s_0, s_1, \dots, s_n])$ from $R(s)$.
- **Additive** rewards:

$$U_r([s_0, s_1, \dots, s_n]) = R(s_0) + R(s_1) + \dots + R(s_n)$$

as above.

- **Discounted** rewards:

$$U_r([s_0, s_1, \dots, s_n]) = R(s_0) + \gamma R(s_1) + \dots + \gamma^n R(s_n)$$

where the **discount factor** γ is a number between 0 and 1.

- The discount factor models the preference of the agent for current over future rewards.

How utilities are calculated

- There is an issue with infinite sequences with additive, undiscounted rewards.
 - What will the utility of a policy be?

How utilities are calculated

- There is an issue with infinite sequences with additive, undiscounted rewards.
 - What will the utility of a policy be?
- Unbounded
- ∞ or $-\infty$.
- This is problematic if we want to compare policies.

How utilities are calculated

- Some solutions are (definitions follow):
 - Proper policies
 - Average reward
 - Discounted rewards

How utilities are calculated

- **Proper policies** always end up in a terminal state eventually.
- Thus they have a finite expected utility.

How utilities are calculated

- We can compute the **average reward** per time step.
- Even for an infinite policy this will (usually) be finite.

How utilities are calculated

- Assume: $0 \leq \gamma < 1$ and rewards are bounded by R_{max}
- With discounted rewards the utility of an infinite sequence is finite:

$$\begin{aligned}U_r([s_0, s_1, \dots, s_n]) &= \sum_{t=0}^n \gamma^t R(s_t) \\&\leq \sum_{t=0}^{\infty} \gamma^t R(s_t) \\&\leq \sum_{t=0}^{\infty} \gamma^t R_{max} \\&\leq \frac{R_{max}}{(1 - \gamma)}\end{aligned}$$

Optimal policies

- With discounted rewards we compare policies by computing their expected values.
- The expected utility of executing π starting in s is given by:

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

where S_t is the state the agent gets to at time t .

- S_t is a random variable and we compute the probability of all its values by looking at all the runs which end up there after t steps.

Optimal policies

- The optimal policy is then:

$$\pi^* = \arg \max_{\pi} U^{\pi}(s)$$

- It turns out that this is independent of the state the agent starts in.

Optimal policies

3	0.812	0.868	0.918	<div>+ 1</div>
2	0.762		0.660	<div>-1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

- Here we have the values of states if the agent executes an optimal policy

$$U^{\pi^*}(s)$$

Optimal policies

3	0.812	0.868	0.918	+ 1
2	0.762		0.660	- 1
1	0.705	0.655	0.611	0.388
	1	2	3	4

- Here we have the values of states if the agent executes an optimal policy

$$U^{\pi^*}(s)$$

- What should the agent do if it is in (3, 1)?

Example

- The answer is *Left*.
- The best action is the one that maximises the expected utility.
- (You have to calculate the expected utility of all the actions to see why Left is the best choice.)