

6CCS30ME/7CCSMOME – Optimisation Methods

Lecture 2

Single-source shortest-paths:

Dijkstra's algorithm, shortest-paths algorithm for DAGs

Tomasz Radzik and Kathleen Steinhöfel

Department of Informatics, King's College London

2020/21, Second term

Topics

- Single-source shortest-paths; restricted cases
 - Only non-negative edge weights allowed:
Dijkstra's shortest-paths algorithm
 - The input graph is acyclic (a DAG – a directed acyclic graph):
Single-source shortest paths algorithm for DAG's
- In both cases, the Bellman-Ford shortest-paths algorithm can be used.

In both cases, the new algorithms are faster than the Bellman-Ford algorithm.

Dijkstra's shortest-paths algorithm

- **Crucial assumption:** all edge weights are nonnegative.
- For convenience, assume that all nodes are reachable from s .

node	a	b	c	h	r	s	v	x	y	z
PARENT	nil	nil	nil	nil	nil	nil	nil	nil	nil	nil
d	∞	∞	∞	∞	∞	0	∞	∞	∞	∞

$\text{DIJKSTRA}(G, w, s)$ $\{ G = (V, E) \}$
 INITIALIZATION(G, s) $\{ \text{"relaxation technique" initialization} \}$
 $S \leftarrow \emptyset$ $\{ \text{nodes } v \text{ for which we know that } d[v] = \delta(s, v) \}$
 $\underline{Q \leftarrow V}$ $\{ \text{the other nodes in } \underline{\text{Priority Queue}} \text{ with keys } d[.] \}$

while $\underline{Q \neq \emptyset}$ **do**


$u \leftarrow \underline{\underline{\text{EXTRACT-MIN}(Q)}}$ $\{ u \text{ has the min } d[.] \text{ value in } Q \}$

$S \leftarrow S \cup \{u\}$

for each node $v \in \text{Adj}[u]$ **do** RELAX(u, v, w)


end_while .

- Q is implemented as a **Priority Queue** data structure.
 Priority Queue maintains pairs (value, key)
 and the main operation is EXTRACT-MIN.
 In Dijkstra's algorithm: pairs (node, $d[\text{node}]$).




Example

From [CLRS] textbook:



Example: initialization (continued at LGT)

From [CLRS] textbook:



$$Q = \{ (s, 0), (u, \infty), (v, \infty), (x, \infty), (y, \infty) \}$$


The correctness of Dijkstra's algorithm

An intermediate state of the computation (at the end of one iteration).

- Set S grows by one node per one iteration

End 1-st iter.: $S = \{s\}$


End last iter.: $S = V$



- All edges from nodes in S , and only these edges, have been relax'ed.
- All nodes in S or at the end of edges from S :
 - have finite $d[.]$ values (by induction)
 - have parents, forming a tree (no negative cycles and Prop. (8)).
- For all other nodes, $d[.] = \infty$, and not in the parent tree.

The correctness of Dijkstra's algorithm (2)

At the end of
the computation:



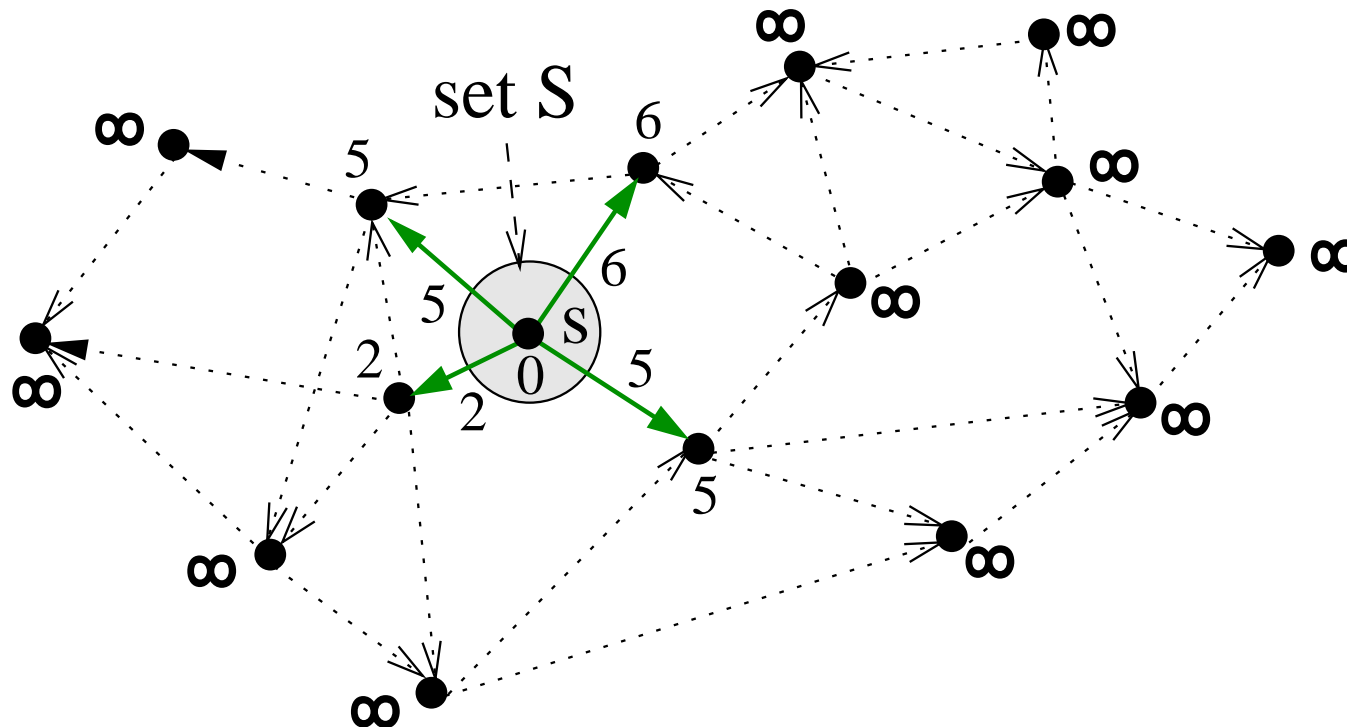
- Set $S = V$.
For each $v \in V$, $d[v]$ is finite and $d[v] \geq \delta(s, v)$.
The parent subgraph is a tree (rooted at s) which includes all nodes.
- We haven't shown yet that the computed:
tree is a shortest paths tree and
the $d[.]$ values are the shortest-path weights.

The correctness of Dijkstra's algorithm (2): the crucial property

The crucial property (the invariant of the computation of Dijkstra's algorithm):

At the end of each iteration (on the main loop) of Dijkstra's algorithm, for each node $v \in S$, $d[v] = \delta(s, v)$.



- This property can be shown **by induction** on the number of iterations.
- **Basis of the induction:** The property holds at the end of the first iteration: $S = \{s\}$, $d[s] = 0 = \delta(s, s)$.



The invariant of Dijkstra's algorithm: the induction step


Induction step:

- Assume the invariant holds at the end of some iteration (not the last one):
for each $v \in S$,
 $d[v] = \delta(s, v)$.
- Let u denote the node selected in the next (current) iteration.
- Show that the invariant holds at the end of the current iteration, (after u added to set S).
- That is, show that
 $d[u] = \delta(s, u)$.



The invariant of Dijkstra's algorithm: the induction step (cont.)

- We have to show that at the beginning of the current iteration, $d[u]$ is the shortest-path weight from s to u . That is, we have to show that $d[u] = \delta(s, u)$.
- We have $d[u] \geq \delta(s, u)$ (relaxation technique). We show $d[u] \leq \delta(s, u)$.
- Take any (simple) path R from s to u and show that its weight $w(R) \geq d[u]$.
- Let z be the first node on path R which is outside S , and let y be the predecessor of z on R (so $y \in S$).
- Let $R1$ be the initial part of path R ending at node y , and let $R2$ be the part of path R from node z to the final node u .
- $$\begin{aligned} w(R) &= w(R1) + w(y, z) + w(R2) \geq d[y] + w(y, z) + w(R2) \geq d[z] + w(R2) \\ &\geq d[u] + w(R2) \geq d[u]. \end{aligned}$$
- The inequalities above follow from: the inductive assumption, RELAX(y, z) done, the rule for selecting u , and the non-negative weights of edges, respectively.
- Thus no path from s to u has smaller weight than $d[u]$, so $d[u] \leq \delta(s, u)$.



The running time of Dijkstra's algorithm

- n – number of nodes; m – number of edges.
- Initialisation (as in the relaxation technique): $\Theta(n)$.
- For every edge (u, v) , $\text{RELAX}(u, v)$ is performed exactly once.
This happens during the iteration when node u is removed from Q .
(Each node is removed from Q exactly once.)
- The running time depends on the implementation of the priority queue Q .
- Note that $n - 1 \leq m \leq n(n - 1)$, so $m = \Omega(n)$ and $m = O(n^2)$.
- The naive implementation of Q , using an unordered list of elements:



- one $\text{EXTRACT-MIN}(Q)$: $O(n)$ time; all of them: $O(n^2)$;
 - one RELAX operation: $O(1)$; all of them: $O(m)$;
 - the total running time: $\Theta(n) + O(n^2) + O(m) = O(n^2)$.
- What would be the running time of Dijkstra's algorithm, if Q was maintained as an ordered list? (LGT).

Priority Queue implemented as Heap: operations and performance

- To improve the running time of Dijkstra's algorithm, use the **Heap** implementation of the Priority-Queue to maintain Q .
- The main Priority Queue operations in Heap (n denotes the size of the heap, that is, the number of elements in the heap):
 - $\text{INSERT}(Q, v, k)$: insert the value-key pair (v, k) . $O(\log n)$ time
 - $\text{EXTRACT-MIN}(Q)$: remove and return the pair with the smallest key. $O(\log n)$ time
 - Check if Heap is not empty. $O(1)$ (constant) time
 - Initialise Heap with given n elements. $\Theta(n)$ time.
- In Dijkstra's algorithm, we also need heap operation:
 $\text{HEAP-DECREASE-KEY}(Q, v, k)$
 - decrease the key associated with v to k . $O(\log n)$ time

Dijkstra's algorithm with Heap

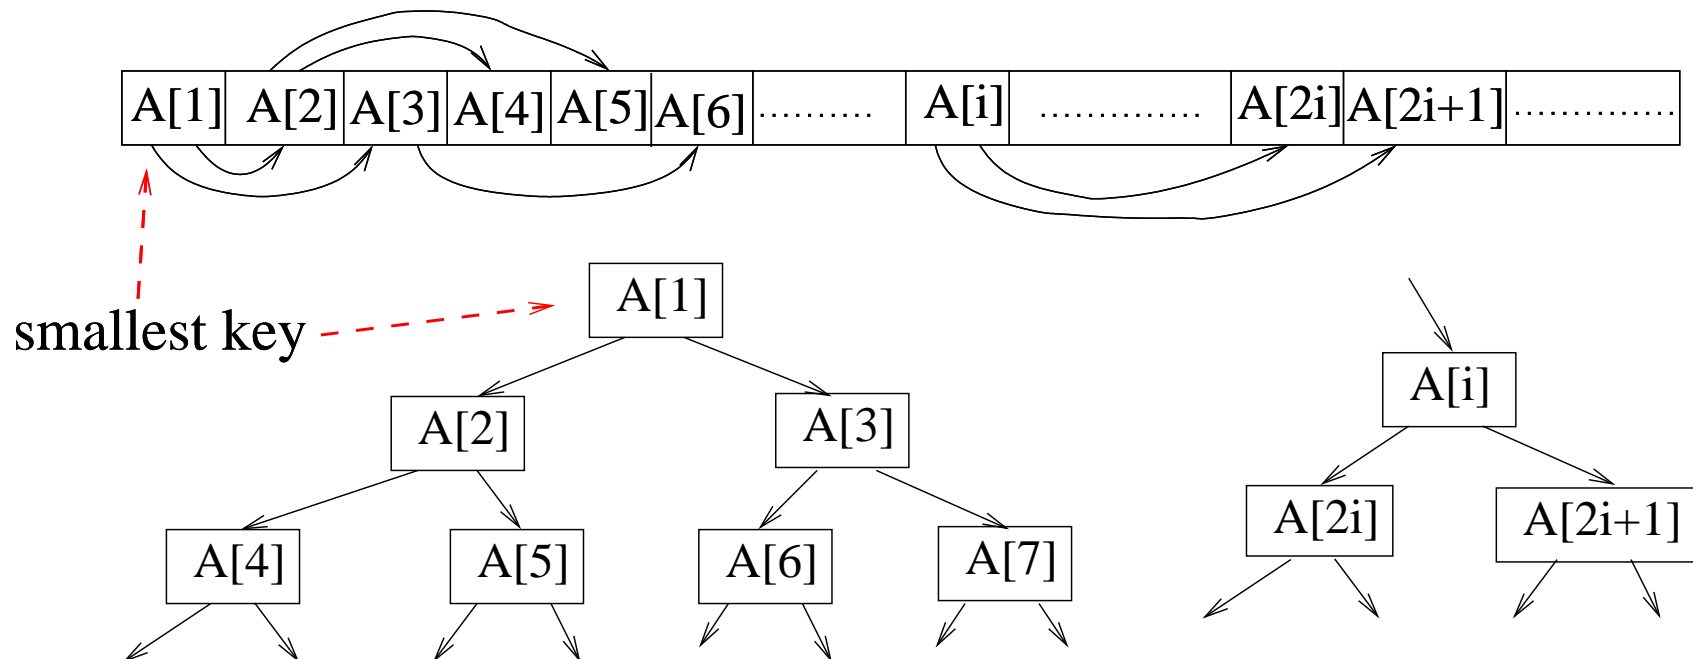
- The set Q in Dijkstra's algorithm maintained as **Heap**:
 - Initialisation of the Heap: $\Theta(n)$ time.
 - One EXTRACT-MIN(Q): $O(\log n)$ time; all: $O(n \log n)$ time.
 - one RELAX(u, v): $O(\log n)$ time, since it may involve one operation HEAP-DECREASE-KEY; all: $O(m \log n)$ time.
 - The total running time of Dijkstra's algorithm with Heap is: $\Theta(n) + \Theta(n) + O(n \log n) + O(m \log n) = O(m \log n)$.
- If the input graph is **dense**, that is (here), if $m = \Omega(n^2 / \log n)$, then the implementation of Dijkstra's algorithm **without Heap** (maintaining Q as an unordered list) gives the better (worst-case) running time – $O(n^2)$.
- For the other cases (when $m = O(n^2 / \log n)$), the implementation of Dijkstra's algorithm **with Heap** gives the better (worst-case) running time.
- In most applications of Dijkstra's algorithm, graphs are not dense, so Dijkstra's algorithm is commonly assumed to use Heap.

Heap implementation of Priority Queue data structure

- **Heap:** An array $A[1..n]$ with a sequence of numbers (keys) and associated data (values), satisfying some specific partial order of the entries. This specific order is called the **heap property** (Min-Heap here):

$$A[i] \leq A[2i] \quad \text{and} \quad A[i] \leq A[2i + 1], \quad \text{for } i = 1, 2, \dots$$

Below, an arrow points from a smaller key to a larger key:




- The operations of extracting the minimum and inserting a new element take $O(\log n)$ time, where n is the current size of the Heap.

Heap in Dijkstra's algorithm

The entries in the Heap array are the pairs $(u, d[u])$, where u is a node in the graph and $d[u]$ is its current shortest path estimate. The number $d[u]$ is the key of this entry.

	1	2	3	4	
Q_A:	x	y	z	v
	d[x]	d[y]	d[z]	d[v]

		y		v		x		z	
pos_in_Q:		2		4		1		3	




- We need a heap operation **HEAP-DECREASE-KEY**, which restores the heap property when the key of one entry is decreased.

RELAX (u, v) (in Dijkstra's algorithm, if Q is a Heap)


if $d[v] > d[u] + w(u, v)$ **then**
 $d[v] \leftarrow d[u] + w(u, v);$ **PARENT** $[v] \leftarrow u$
HEAP-DECREASE-KEY $(Q, v, d[v])$

Operation **HEAP-DECREASE-KEY** takes $O(\log n)$ time.

Heap Extract-Min operation




Heap-Decrease-Key operation



Single-source shortest paths in DAG's

Input:

- $G = (V, E)$ – directed acyclic graph (DAG);
- w – weights of edges (may be negative);
- $s \in V$ – the source node.



Output: the shortest-path weights from s to all other nodes and a shortest-paths tree from s .

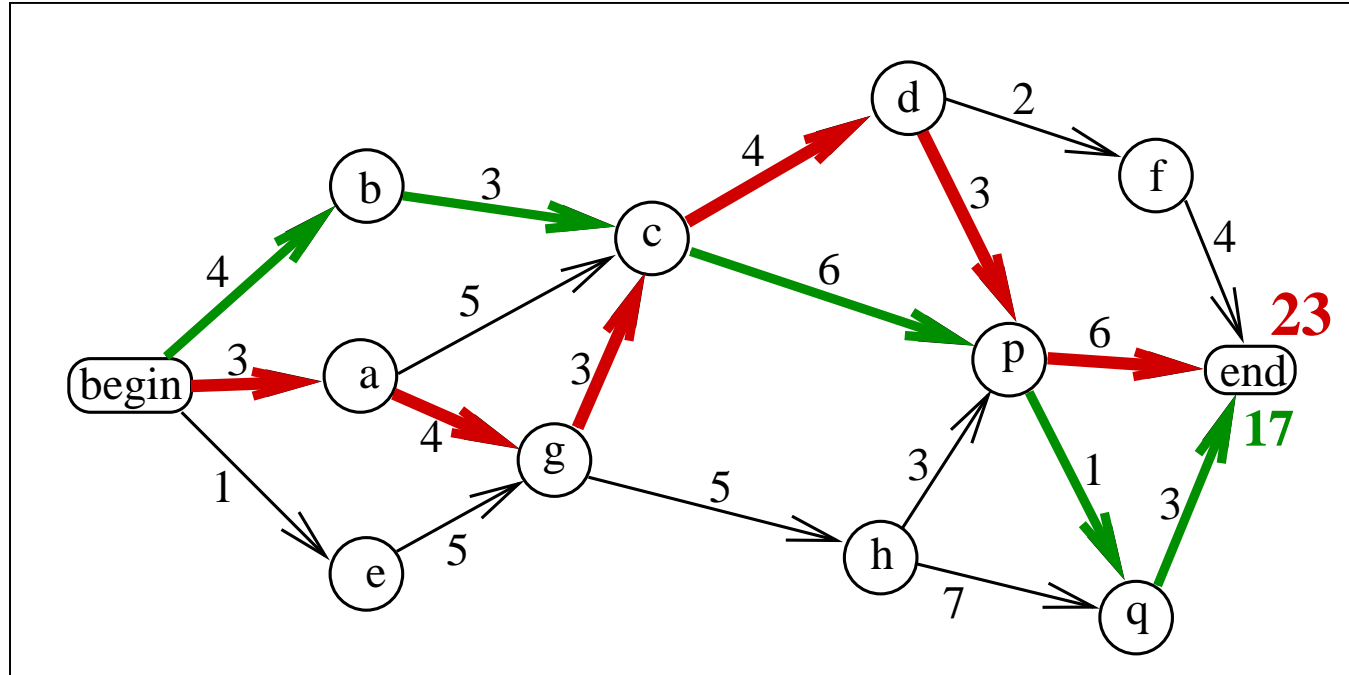
There may be edges with negative weights, but no problem with negative cycles, because there are no cycles at all.

Bellman-Ford: $O(mn)$.

We show an algorithm which runs in $O(n + m)$ time.
(Linear, optimal time)


Application

- Determine the critical (longest) paths in PERT chart analysis (Program Evaluation and Review Technique).
- Nodes: milestones of the project ('synchronisation points') . Edges: tasks of the project. The weight of an edge: the duration of this task.
Find the longest path from node "begin" to node "end":
this gives the fastest possible time for completing the whole project (that is, for completing all tasks of the project).




DAGs: Longest paths to shortest paths by negating all edge weights

- To compute longest paths from the start node, negate all edge weights and compute shortest paths from the start node.



Algorithm

- Consider the nodes in a topological order: all edges go in the same direction.



DAGSHORTESTPATHS(G, w, s)

- 1: topologically sort the nodes of G { put the nodes in a topological order }
- 2: INITIALIZATION(G, s)
- 3: **for** each node u taken in the topological order computed in step 1 **do**
- 4: **for** each node $v \in Adj[u]$ **do**
- 5: RELAX(u, v, w)

- Running time: $O(n + m)$
(topological sort takes $O(n + m)$ time).
- Correctness: show (by induction) that when node u is considered in line 3, then $d[u] = \delta(s, u)$.

Examples and Exercises – LGT

1. Example of the computation of Dijkstra's algorithm – slides 4-5.
2. What would be the running time of Dijkstra's algorithm, if the priority queue Q was maintained as an ordered list?


Examples and Exercises – LGT (cont)

3. (Exercise 24.2-3 from [CLRS])

The PERT chart formulation given in this lecture is somewhat unnatural. More naturally, vertices would represent tasks and edges would represent order constraints; edge (u, v) would indicate that task u must be performed before task v . We would then assign weights (duration of the tasks) to vertices, not edges. Modify the DAGSHORTESTPATHS algorithm so that it finds, in linear time, a longest path in a directed acyclic graph with weighted vertices. Let $w(v)$ denote the weight of a vertex v . Trace the computation of the algorithm on the graph below.

The green path
has weight 20.


The red path,
with weight 26,
is the longest path.



Exercises – SGT

1. This exercise shows that Dijkstra's algorithm does not necessarily compute shortest paths, if there are negative weights; even if there are no negative cycles.

- (a) Show the values $d[x]$ and the tree computed by Dijkstra's algorithm for the input graph:



- (b) Show the shortest-path weights and a shortest-paths tree in this graph.
- (c) How to modify Dijkstra's algorithm so that it runs also for graphs where edge weights may be negative?
- (d) What is the running time of the modified algorithm?

Exercises – SGT (cont.)

2. Design a linear-time ($O(n + m)$ -time) algorithm which for a given directed acyclic graph (with n nodes and m edges) computes a topological order of the nodes.

Do not use the Depth-First Search (DFS) algorithm.

(There is a topological sort algorithm which is based on the DFS algorithm, but in this exercise you are asked to design an alternative algorithm.)

Hint: How to identify the first node for the topological order? How to identify subsequent nodes?

For the running time, assume the adjacency-list representation of the graph.

Specify all data structures which your algorithm needs to achieve the $O(n + m)$ running time.


Exercises – SGT (cont.)

3. Revise the Breadth-First Search (BFS) algorithm.

Trace the execution of BFS on the graph shown below.


4. Revise the Depth-First Search (DFS) algorithm.

Trace the execution of BFS on the graph shown below.



Exercises – SGT (cont.)

5. The graph below has a negative cycle reachable from the source. Trace the computation of the Bellman-Ford algorithm with FIFO queue on this graph until the PARENT edges form a cycle.



Assume that the edges outgoing from nodes are given in this order:

$(s, u), (s, x);$
 $(u, x), (u, v);$
 $(x, u), (x, y);$
 $(v, x), (v, y);$
 $(y, s), (y, v).$